

Interests: Data-Centric/Trustworthy AI & Decisions, Reinforcement Learning, Probability, Optimization

Education

- **The Chinese University of Hong Kong, Shenzhen, China** 2025 (expected)
Ph.D. in Computer Science and Information Engineering.
Supervisor: [Zhi-Quan \(Tom\) Luo](https://en.wikipedia.org/wiki/Zhi-Quan_Tom_Luo) https://en.wikipedia.org/wiki/Zhi-Quan_Tom_Luo
Thesis Committee: *Hongyuan Zha, Xinyun Chen, Baoxiang Wang, John C.S. Lui, Benjamin Van Roy (Stanford & Google DeepMind)*
- **Huazhong University of Science and Technology, Wuhan, China** 2020
M.S. in Computer Science.
Ranked 1st/134 overall, 1st/26 in Computer Theory and Software specialization.
- **Huazhong University of Science and Technology, Wuhan, China** 2017
B.E. in Computer Science (Honors Program). Outstanding Graduate.

Professional Experience

- **Research Institute of Big Data, Shenzhen, China** 2023 - present
Research Assistant
 - Developed **learning and optimization techniques in strategic games** for sensing and transportation systems, focusing on leveraging payoff correlation to **mitigate the curse of multi-agent**.
 - Contributed to the HuatuoGPT project, a **medical decision-making language model agent deployed in 16 hospitals**: (1) Data selection and synthesis for multi-turn medical conversation and decision-making. (2) Multi-turn actor-critic type LLM agent in outpatient referral scenarios.
- **Tencent AI & Robotics X, Shenzhen, China** 2019 - 2022
Research Intern, Agent Center
Data-efficient Reinforcement Learning (RL):
 - Contributed to a high-throughput distributed Actor-Learner system, enabling parallel on-policy roll-out and stable policy optimization with off-policy data reuse. Research published in NeurIPS.
 - Led a long-term research project on scalable exploration for RL in complex environments via scalable uncertainty estimation. Developed **HyperAgent**, published in ICML, achieving 7x data efficiency and 20x computation reduction with theoretical guarantees.
- **SenseTime Research, Peking, China** 2018
Computer Vision Trainee Researcher
Worked on an incremental and continual learning system for multi-label image classification.
- **Department of Computer Science, Cornell University, Ithaca, NY** 2017
Undergraduate Research Assistant
Worked with John E. Hopcroft on hidden community detection, a novel graph-theoretical concept.
- **Microsoft Research Lab, Asia** 2016
Research Intern in Theory Center
Worked on influence learning and maximization in social networks.
- **John Hopcroft Center on Computing Science, Wuhan, China** 2015-2017
Undergraduate Research Assistant
Led a reading group on machine learning and spectral graph theory. Worked on network analysis.

Awards

- **Best Paper Award**, in the 3rd doctoral and postdoctoral Daoyuan academic forum, 2024.
- **Best Student Paper**, in IEEE Sensor Array and Multichannel Signal Processing Workshop, 2024.

Awards (continued)

- **SRIBD Ph.D. Fellowship** (Gold Class), by Shenzhen Research Institute of Big Data (SRIBD), 2023.
- **Presidential Ph.D. Fellowship**, by The Chinese University of Hong Kong, Shenzhen, 2019–2023.
- **Tencent AI Ph.D. Fellowship**, by Tencent & The Chinese University of Hong Kong, Shenzhen, 2018.
- **National Scholarship**, by Huazhong University of Science and Technology, 2018.
- **Qiming Star Award (Selected as one of 5 recipients out of 7,112 undergraduates.)**, by Huazhong University of Science and Technology, 2016. **Reports:** [1] [Newspaper](#). [2] [HUST Online](#).
- **Award of Excellence in Internship**, by Microsoft Research Lab, 2016.
- **First Prize**, in Parallel computation and Application Contest (PAC) held by Intel and CCF, 2015.
- **First Prize**, in China National Mathematics Olympiad (Province-level Math League), 2012.

Selected Oral Presentations

- **GPT-HyperAgent: Adaptive Foundation Models for Online Decisions**
Invited talk in 2024 INFORMS Annual Meeting, Seattle, Oct. 21, 2024.
- **HyperAgent: Advancing Scalable Exploration through Fast Uncertainty Estimation in RL**
Invited talk in International Symposium on Mathematical Programming (ISMP), Montréal, Jul. 25, 2024.
- **HyperAgent: A Simple, Efficient and Scalable RL Framework for Complex Environments**
a.k.a. "Q* meets Thompson Sampling: Scaling up Exploration via HyperAgent"
a.k.a. "Scaling Uncertainty Estimation for Online Decisions"
Invited talk at MIT, Jul. 30, 2024.
Invited talk at RLChina.org, Jun. 25, 2024.
Invited talk at Princeton University, May 2, 2024.
Invited talk in INFORMS Optimization Society (IOS) Conference, Rice University, Mar. 23, 2024.
Contributed talk, in the third doctoral and postdoctoral Daoyuan academic forum, Jan. 13, 2024.
- **No-Regret Learning in Unknown Game with Applications**
Invited talk in RL Theory Student Workshop at Nanjing University, Aug. 23, 2022.
Contributed Talk in the second doctoral and postdoctoral Daoyuan academic forum, Aug. 20, 2022.
- **HyperDQN: A Randomized Exploration Method for Deep Reinforcement Learning**
Contributed Talk in NeurIPS Workshop Ecological Theory of Reinforcement Learning, Dec. 14, 2021

Selected Research Publications

Preprints

- 1 **Yingru Li**. *Simple, Unified Analysis of Johnson-Lindenstrauss with Applications*. Preprint. Presentation at ICML 2024 Workshop "High-dimensional Learning Dynamics 2024: The Emergence of Structure and Reasoning". arXiv: [2402.10232](#) [stat.ML].
- 2 **Yingru Li**, Liangqi Liu, Hao Liang, Wenqiang Pu, and Zhi-Quan Luo. *Optimistic Thompson Sampling for No-Regret Learning in Unknown Games*. Preprint. Presentation at ICML 2023 Workshop "The Many Facets of Preference-Based Learning". arXiv: [2402.09456](#) [cs.LG].
- 3 **Yingru Li**, Jiawei Xu, and Zhi-Quan Luo. *Adaptive Foundation Models for Online Decisions: HyperAgent with Fast Incremental Uncertainty Estimation*. Preprint. Presentation at ICML 2024 Workshops: (1) "Aligning Reinforcement Learning Experimentalists and Theorists"; (2) "Automated Reinforcement Learning: Exploring Meta-Learning, AutoML, and LLMs". arXiv: [2407.13195](#) [cs.LG].
- 4 **Yingru Li**. *Probability Tools for Sequential Random Projection*. Preprint. Presentation at ICML 2024 Workshop "High-dimensional Learning Dynamics 2024: The Emergence of Structure and Reasoning". 2024. arXiv: [2402.14026](#) [math.ST].

Conference Proceedings

- 5 **Yingru Li** and Zhi-Quan Luo. “Prior-dependent analysis of posterior sampling reinforcement learning with function approximation”. In: *The 27th International Conference on Artificial Intelligence and Statistics (AISTATS)*. 2024. arXiv: [2403.11175 \[stat.ML\]](#).
- 6 **Yingru Li**, Jiawei Xu, Lei Han, and Zhi-Quan Luo. “Q-Star Meets Scalable Posterior Sampling: Bridging Theory and Practice via HyperAgent”. In: *The 41st International Conference on Machine Learning (ICML)*. Proceedings of Machine Learning Research. 2024. arXiv: [2402.10228 \[cs.LG\]](#).
- 7 Liangqi Liu, Wenqiang Pu, **Yingru Li**, Bo Jiu, and Zhi-Quan Luo. “Radar Anti-jamming Strategy Learning via Domain-knowledge Enhanced Online Convex Optimization”. In: *2024 IEEE 13th Sensor Array and Multichannel Signal Processing Workshop (SAM)*. IEEE. 2024. arXiv: [2402.16274 \[eess.SP\]](#).
- 8 Ziniu Li, **Yingru Li**, Yushun Zhang, Tong Zhang, and Zhi-Quan Luo. “HyperDQN: A Randomized Exploration Method for Deep Reinforcement Learning”. In: *International Conference on Learning Representations (ICLR)*. 2022. [URL: https://openreview.net/pdf?id=X0nrKAXu7g-](#).
- 9 Qing Wang, **Yingru Li**, Jiechao Xiong, and Tong Zhang. “Divergence-Augmented Policy Optimization”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. Vol. 32. Curran Associates, Inc., 2019. [URL: https://openreview.net/pdf?id=rylacVSeIS](#).

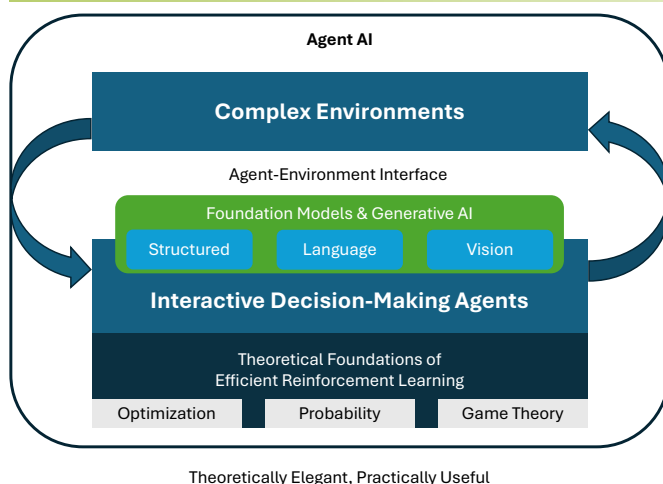
Journal Articles

- 10 Kun He, **Yingru Li**, Sucheta Soundarajan, and John E. Hopcroft. “Hidden community detection in social networks”. In: *Information Sciences* 425 (2018), pp. 92–106. ISSN: 0020-0255. [DOI: https://doi.org/10.1016/j.ins.2017.10.019](#).

Under Preparation

- 11 **Yingru Li**, Yanchuan Tang, Xiaoxiao Liu, Benyou Wang, and Zhi-Quan Luo. “Multi-turn Actor-critic Language Agents for Hospital Outpatient Referral”. 2024.
- 12 Liangqi Liu, Wenqiang Pu, **Yingru Li**, Bo Jiu, and Zhi-Quan Luo. “Learning an Opponent-aware Anti-jamming Strategy via Online Convex Optimization”. 2024.
- 13 Fei Yu, **Yingru Li**, Benyou Wang, and Zhi-Quan Luo. “Uncertainty-guided Multi-step Reasoning in LLMs”. 2024.

Research Statement



My research vision is to develop interactive AI agents that operate reliably in complex, uncertain, dynamic and human-in-the-loop environments. This work necessitates advancements in knowledge and uncertainty

representation, exploration, continual adaptation, reasoning and decision-making. To achieve these goals, I use and develop **fundamental tools** in probability, game theory, and information theory. Meanwhile, I design **scalable algorithms** that take advantages of modern large-scale computation tools.

My methods have been applied to llm math reasoning, strategic games in sensing & transportation, human-AI interplay for content moderation, and reliable medical decision-making. The significance of my work has been recognized through presentations at top-tier conferences (ICML, NeurIPS, ICLR, AISTATS, ISMP, INFORMS), as well as through awards (Best Paper Award at the 2024 Daoyuan Forum, Best Student Paper Award at the 2024 IEEE SAM). For more information, visit <https://richardli.xyz/#research>.

Academic Service

- **Reviewer** for Conference on Neural Information Processing Systems (NeurIPS) [12 papers], International Conference on Learning Representations (ICLR) [2 papers]; ICLR Workshop "Bridging the Gap Between Practice and Theory in Deep Learning" (2 papers), ICML Workshop "Aligning Reinforcement Learning Experimentalists and Theorists" (2 papers).
- **Chair** for [RL Seminar](#) in The Chinese University of Hong Kong, Shenzhen, China (Spring 2019, Summer 2020, Fall 2020, Spring 2021, Summer 2021, Fall 2021, Spring 2022, Fall 2022.); for 2 sessions in INFORMS Annual Meeting 2024 on "Integrating Generative AI with Sequential Decision-making".

Teaching

- | | |
|---|-------------|
| ■ Stochastic Processes (STA/DDA4001) | Fall 2018 |
| ■ Optimization II (MAT3220) | Spring 2019 |
| ■ Distributed and Parallel Computing (CSC4005) | Fall 2019 |
| ■ Reinforcement Learning (DDA6105/CIE6023) | Fall 2020 |
| ■ Matrix Analysis (CIE6002) | Spring 2021 |
| ■ Deep Learning and Their Applications (MDS6224) | Spring 2022 |

My teaching duties include delivering weekly tutorials, correcting assignments, and running laboratory sessions when required, **all in English**.

Beyond Academia

I enjoy photography. I often play tennis and swim, and occasionally play golf. These activities allow me to live in the moment and help me find physical and spiritual freedom.